

Container versus operating system differences

Container:

- niedriger Hardware-Footprint
- isolierte Umgebung
- schnelle Bereitstellung
- Bereitstellung mit mehreren Umgebungen
- Wiederverwendbar

```
/
├── bin
├── etc
├── lib
├── opt
├── proc
├── usr
├── var
│   └── lib
│       └── container
│           └── storage
```

```
/
systemd├──agetty
        ├──auditd──{auditd}
        ├──chronyd
        ├──crond
        ├──dbus-daemon
        ├──irqbalance
        ├──ksmtuned──sleep
        ├──polkitd──6*[{polkitd}]
        ├──rpc.idmapd
        ├──rpc.mountd
        ├──rpc.statd
        ├──rpcbind
        ├──rsyslogd──2*[{rsyslogd}]
        ├──sshd──sshd──bash
        ├──systemd-journal
        ├──systemd-logind
        ├──systemd-machine
        ├──systemd-udevd
        └──< app >
```

Container: namespace mnt, pid, net, user, uts, ipc

```
/
├── bin
├── etc
├── lib
├── opt
├── proc
├── usr
└── var
```

```
/
└── < app > 8080
```

ip: 10.88.0.x

cgroups : mem, cpu

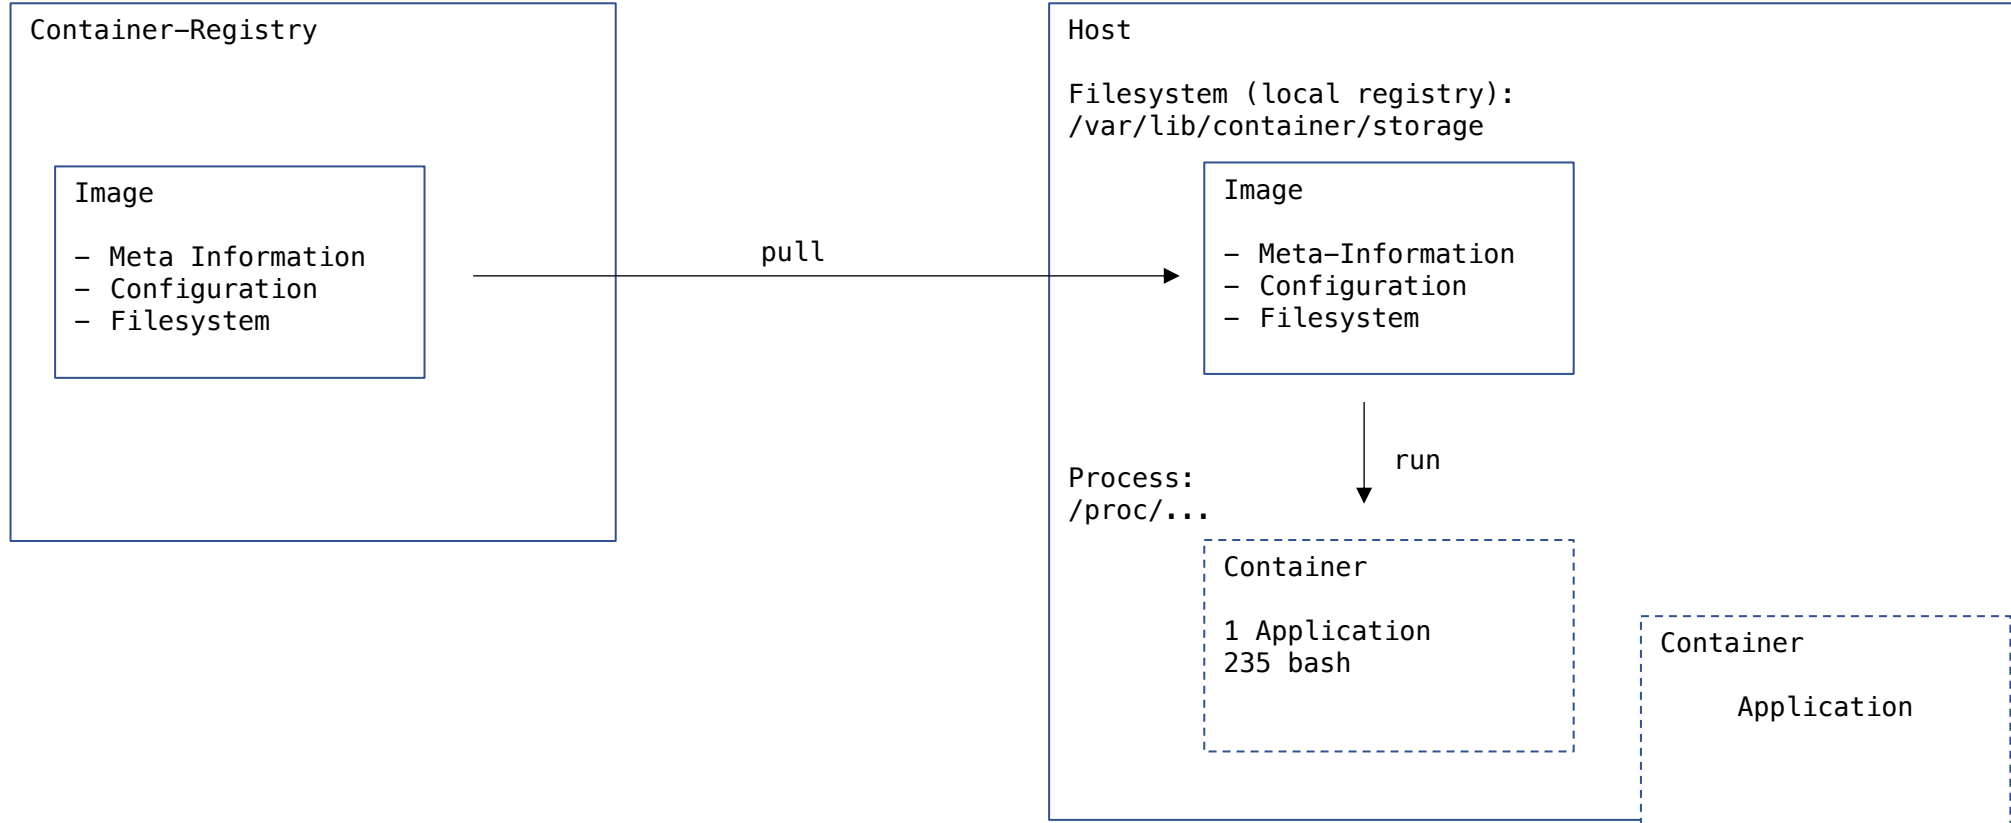
selinux : context security

secomp : syscalls

cni0: 10.88.0.1

eth0

KERNEL



<https://access.redhat.com/RegistryAuthentication>

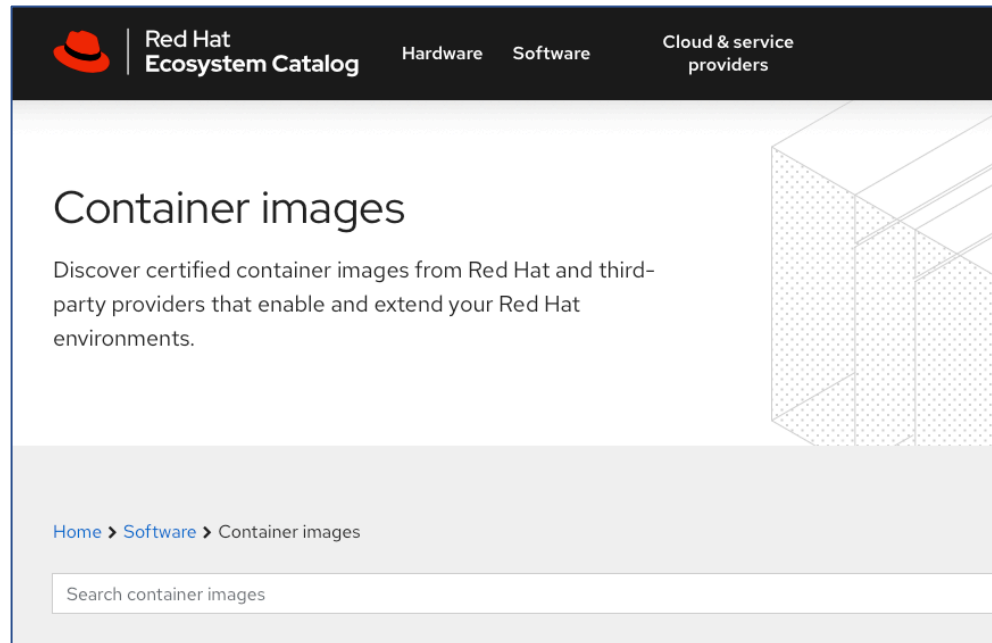
Red Hat Registries

Red Hat distributes container images through three different container registries:

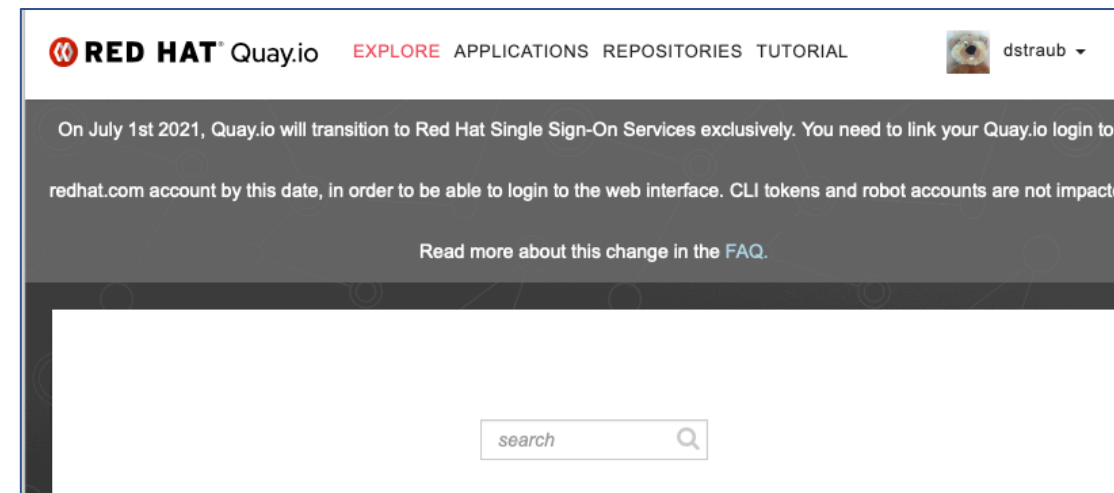
Registry	Content	Supports unauthenticated access	Supports Red Hat login	Supports registry tokens
registry.access.redhat.com	Red Hat products	Yes	No	No
registry.redhat.io	Red Hat products	No	Yes	Yes
registry.connect.redhat.com	Third-party products	No	Yes	Yes

Although both `registry.access.redhat.com` and `registry.redhat.io` hold essentially the same container images, some images that require a subscription are only available from `registry.redhat.io`.

<https://catalog.redhat.com/software/containers/explore>



<https://quay.io>



<https://podman.io>



- Image- und Containermanagement
- OCI: Open Container Initiative
- keine Client/Serverarchitektur
- gleiche Befehlssyntax wie do...
- Kubernetes kompatibel
- `yum install podman`

Open Container Initiative

containers/image

runc

containers/storage

cni

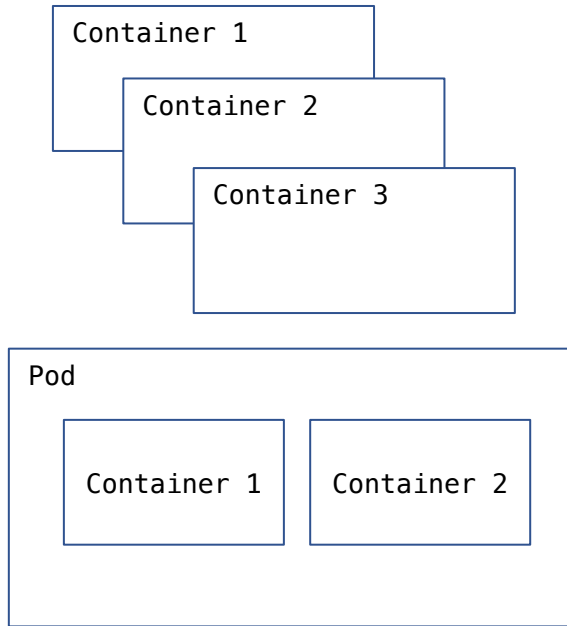
<https://buildah.io>



- Erstellen von Images
- `yum install buildah`

Podman :

großer Aufwand beim Betrieb mehrerer Container, Service-Kommunikation, Routing



Kubernetes : Orchestrierung von Container-Anwendungen

- Service Discovery, Loadbalancing
- Horizontale Skalierung
- Health Checks
- Rolling Updates
- Secret/Configmanagement
- Operatoren: native Kubernetes Anwendungen zum Cluster- und Anwendungs-Management



Openshift (RHOCP):

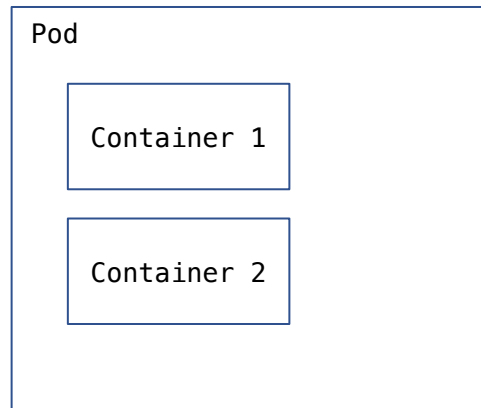
- basiert auf Kubernetes
- Entwickler-Workflow (CI/CD)
- Routing
- Metriken und Log-Management
- einheitliche Benutzeroberfläche

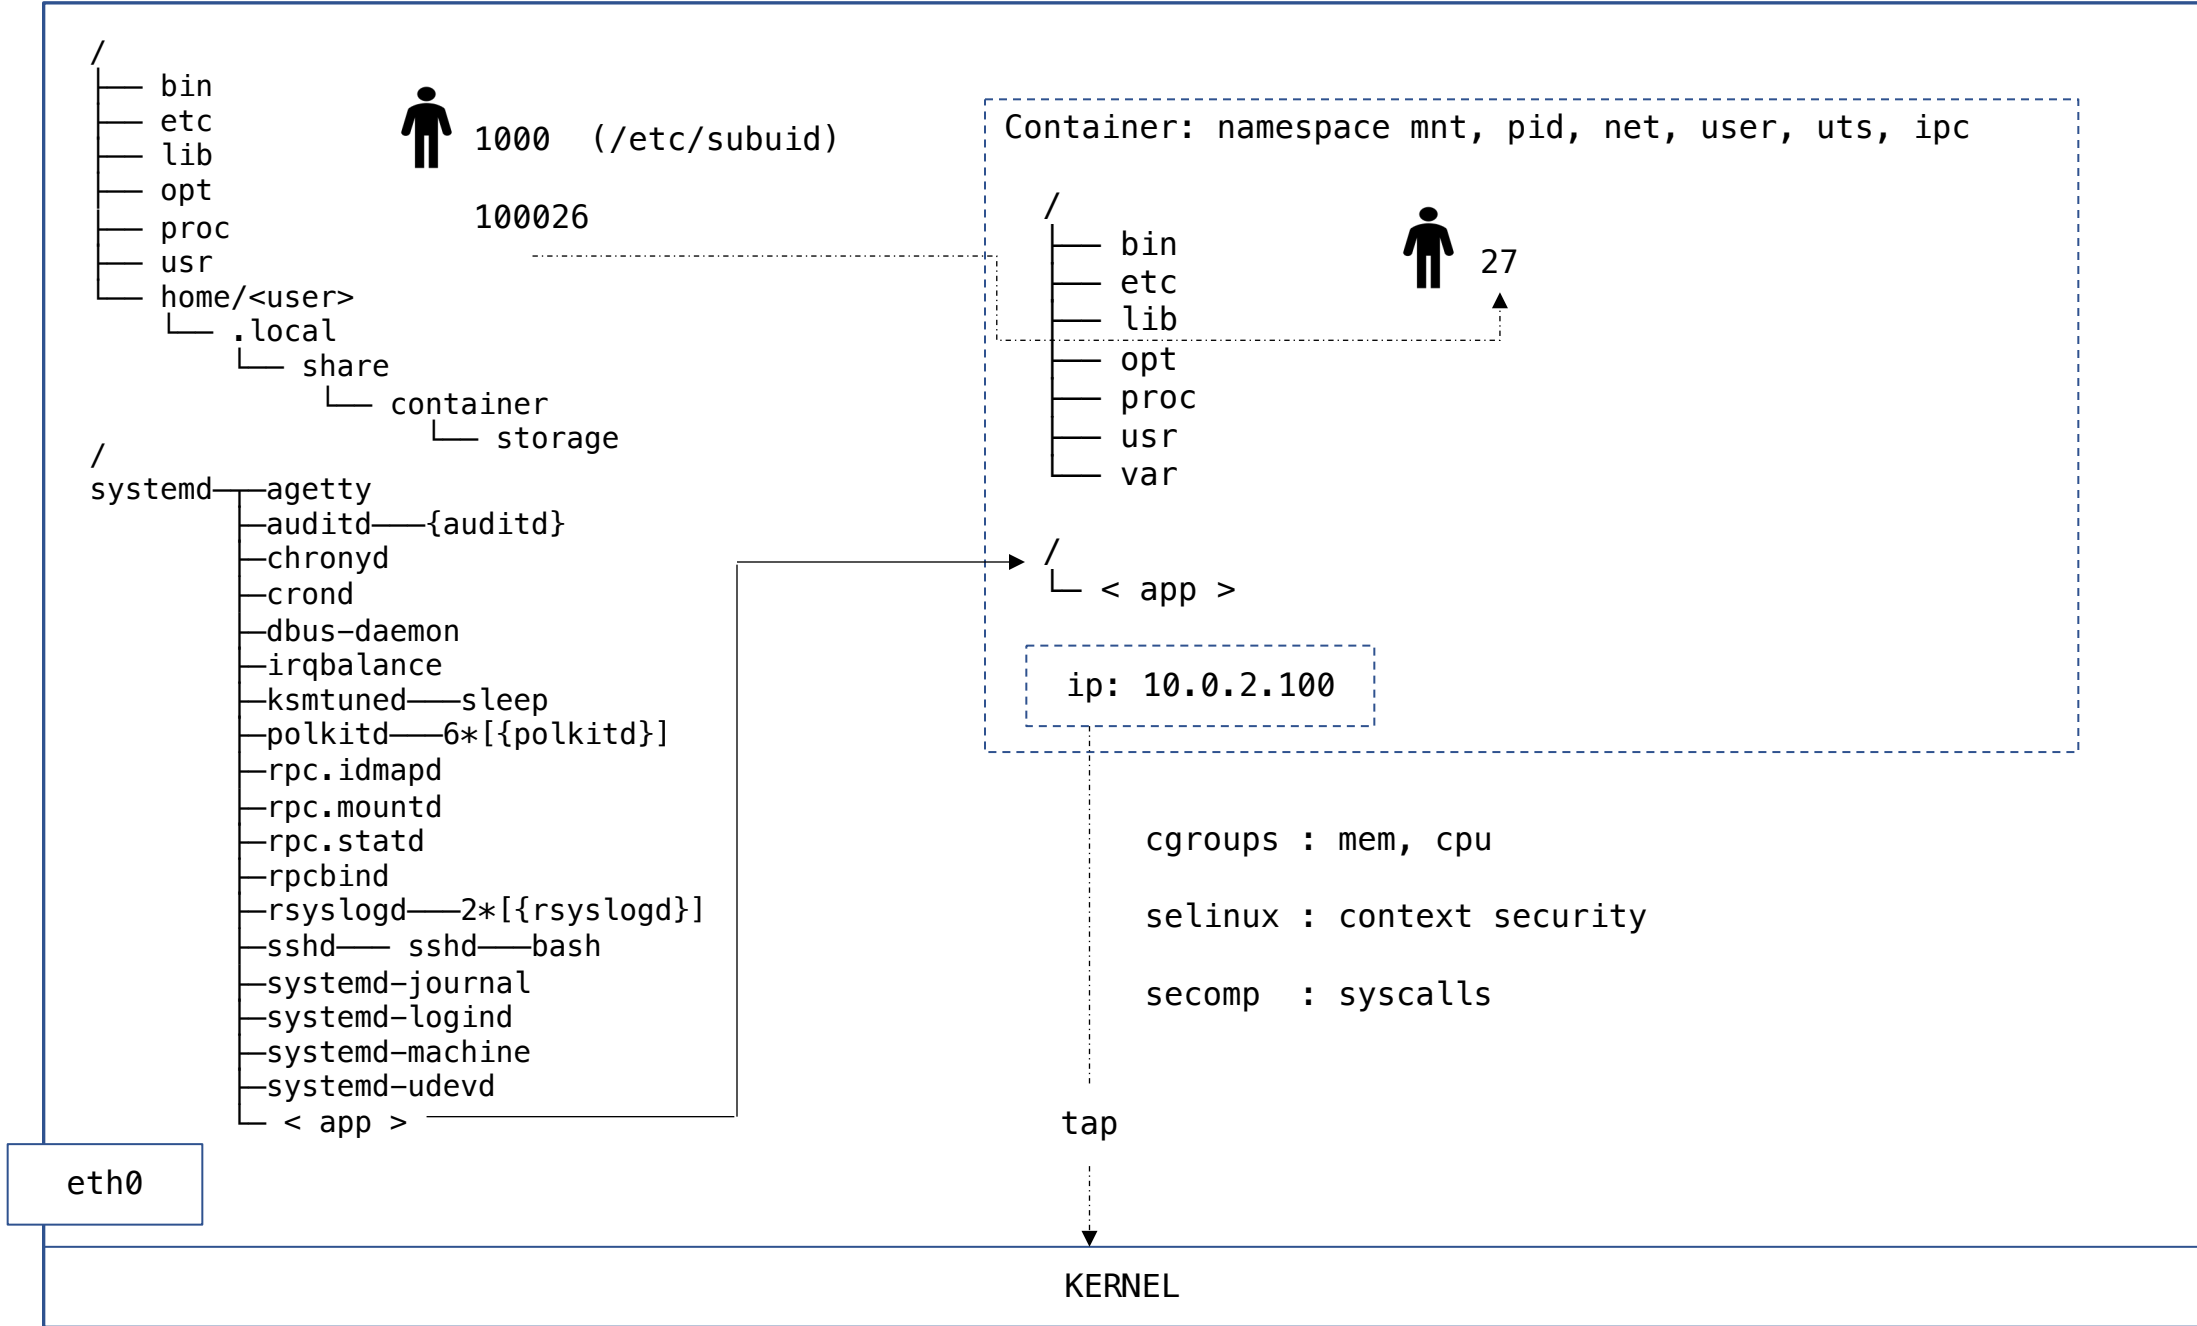
Podman:

- Verwalten von Images und Containern
- mehrere Container können in einen Pod zusammengefasst werden

Kubernetes:

- kleinste Einheit ist der Pod – Gruppe von (unterschiedlichen) Containern
- meistens 1:1 Beziehung (1 Pod enthält ein Container)



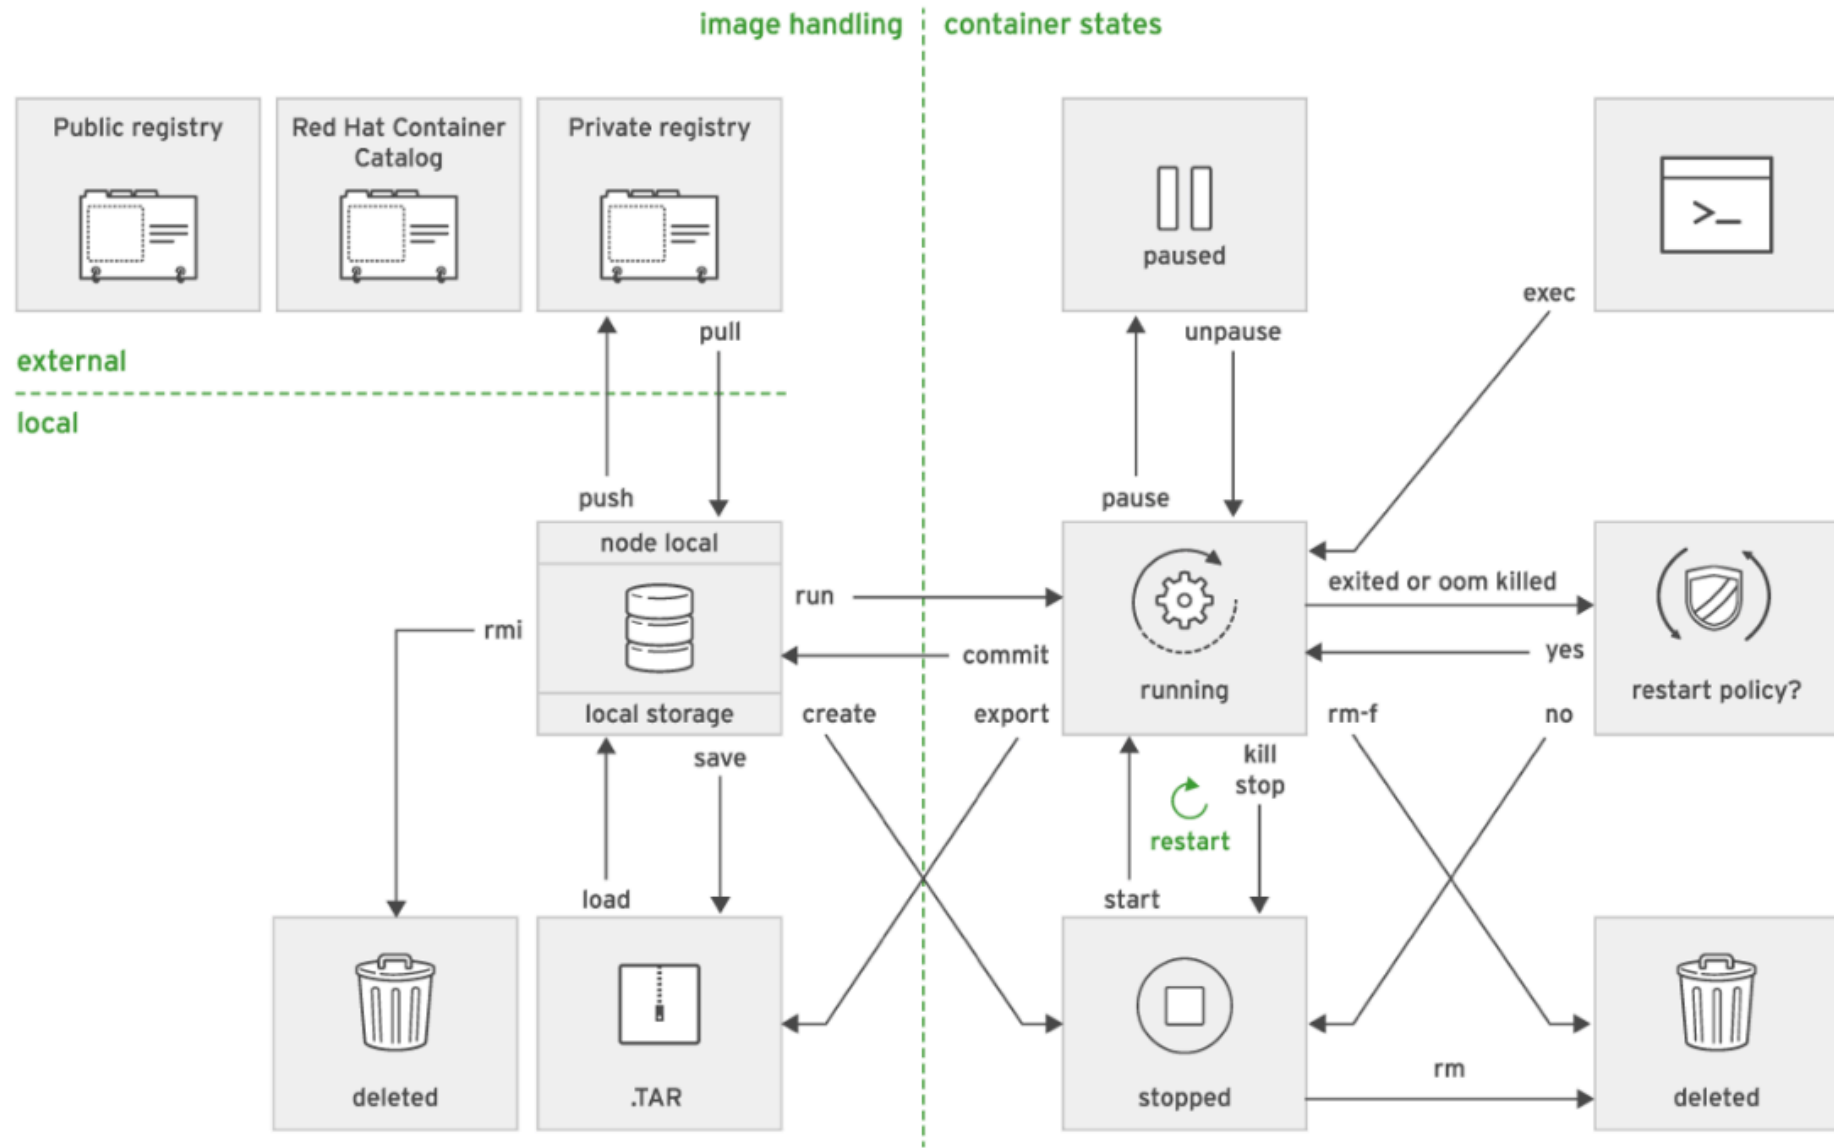


Rootless Container (Linux Kernel > v4.18.0)

- User-Mapping : /etc/subuid , /etc/subgid
 student:100000:65536
 (Container-Benutzer root = Host User)
- Fuse – Filesystem statt Overlay2 (~/.local/share/containers)
- TAP – Network Device (keine reale IP-Adresse)

```
tap0: flags=67<UP,BROADCAST,RUNNING> mtu 65520  
      inet 10.0.2.100 netmask 255.255.255.0 broadcast 10.0.2.255  
      inet6 fe80::6093:deff:febe:f21c prefixlen 64 scopeid 0x20<link>  
      ether 62:93:de:be:f2:1c txqueuelen 1000 (Ethernet)
```

https://github.com/containers/podman/blob/main/docs/tutorials/rootless_tutorial.md



Podman managing subcommands

```
# podman run -d --name httpd rhsc1/httpd-24-rhel7:2.4-36.8
```

```
# podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
4f9e8519685f	.../rhsc1/httpd-24-rhel7:2.4-36.8	/usr/bin/run-http...	About a minute ago	Up (14 seconds ago)		httpd

```
# podman exec httpd cat /etc/hosts
```

```
...  
172.25.250.9    workstation.lab.example.com workstation  
172.25.254.254 classroom.example.com classroom  
172.25.250.254 bastion.lab.example.com bastion  
10.88.0.18    4f9e8519685f
```

```
# curl -I 10.88.0.18:8080
```

```
# podman pause httpd
```

```
# podman ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
4f9e8519685f	.../rhsc1/httpd-24-rhel7:2.4-36.8	/usr/bin/run-http...	2 minutes ago	Paused		httpd

```
# podman unpause httpd
```

```
# podman kill httpd
```

```
# podman logs httpd
```

```
...  
[Mon May 17 17:23:42.147898 2021] [lbmethod_heartbeat:notice] [pid 1] AH02282: No slotmem from mod_heartbeat  
[Mon May 17 17:23:42.153159 2021] [mpm_prefork:notice] [pid 1] AH00163: Apache/2.4.25 (Red Hat) ... resuming normal operations  
[Mon May 17 17:23:42.153196 2021] [core:notice] [pid 1] AH00094: Command line: 'httpd -D FOREGROUND'
```

```
# podman rm httpd
```

```
podman stop: sends SIGTERM, [wait -timeout], send SIGKILL  
podman kill: sends SIGKILL  
podman rm -f -> KILL + rm
```

podman run : Environment

```
podman run -e <KEY>=<VALUE>
```

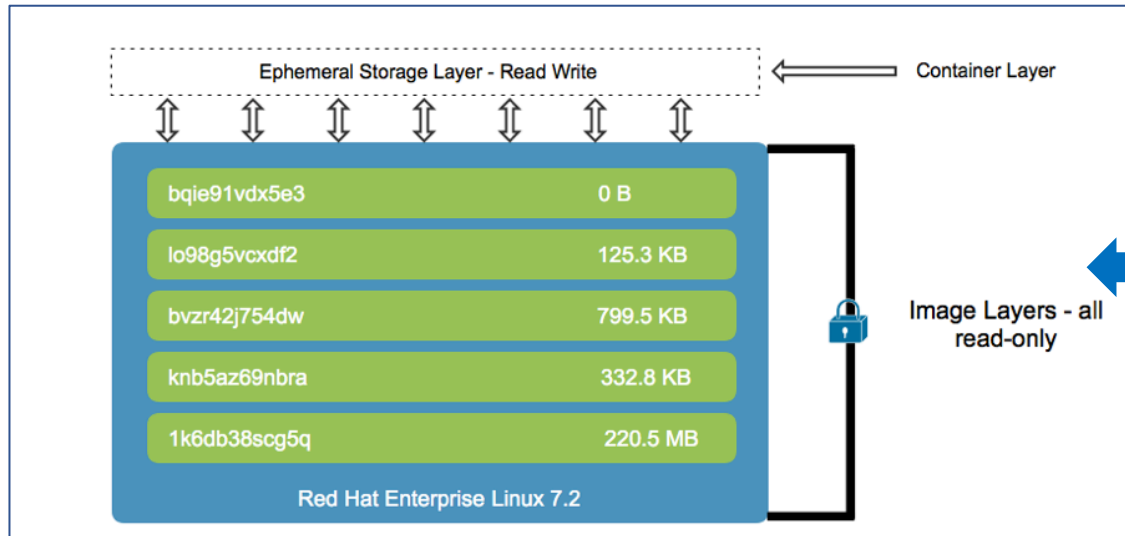
```
podman run --env-file=<host-file>
```

```
podman run --env-host=true|false
```

podman run : Volumes

```
podman run -v <host-dir>:<container-dir>
```

```
podman run --volumes-from <container-name>
```



Mapping (Mount) des Host-Filesystem in den Container:

- Permissions
`chown -R <container-userid> <host-dir>`
`rootless: podman unshare chown -R <container-userid> <host-dir>`
- SELinux
`semange fcontext -a -t container_file_t '<host-dir>(/.*)?'`
`restorecon -Rv <dir>`

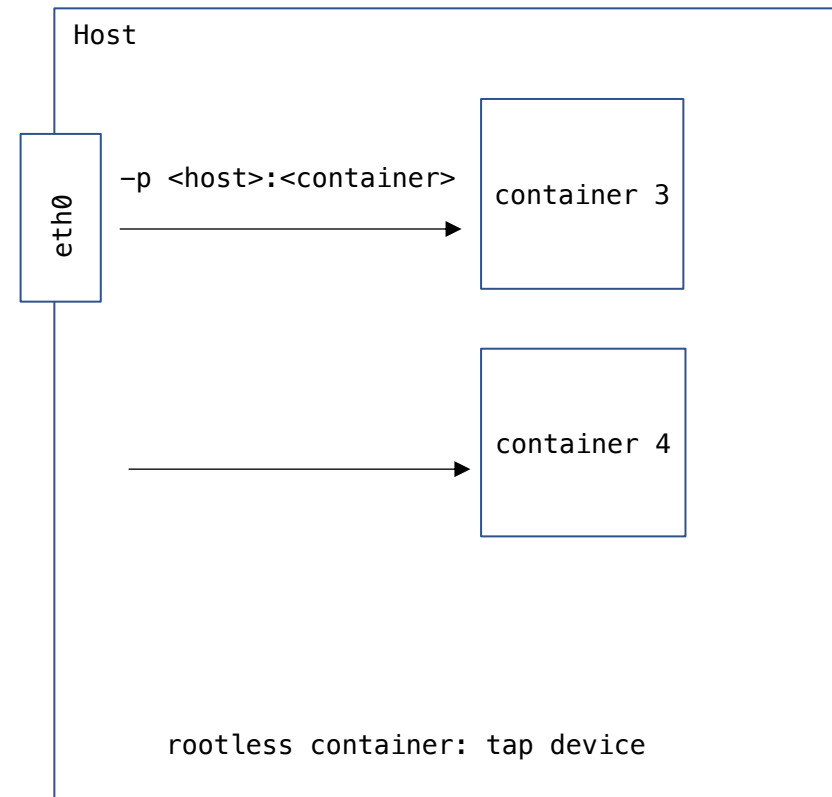
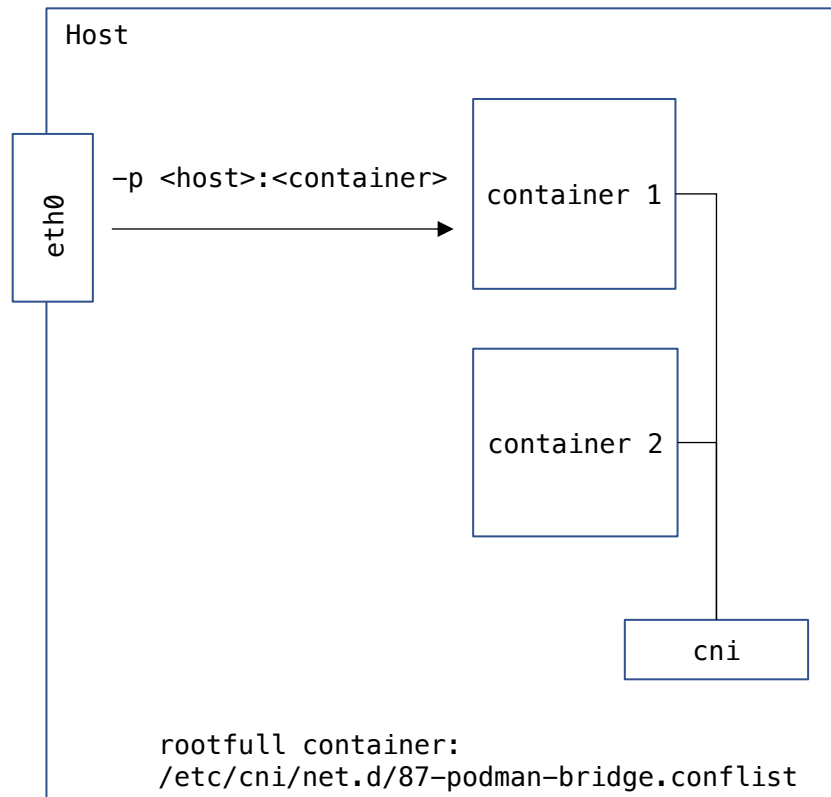
```
podman run -v <host-dir>:<container-dir> <image>
```

podman run - Publishing:

```
podman run -p <host-port>:<container-port> ...
```

```
podman run -P / --publish-all
```

```
podman port -l
```



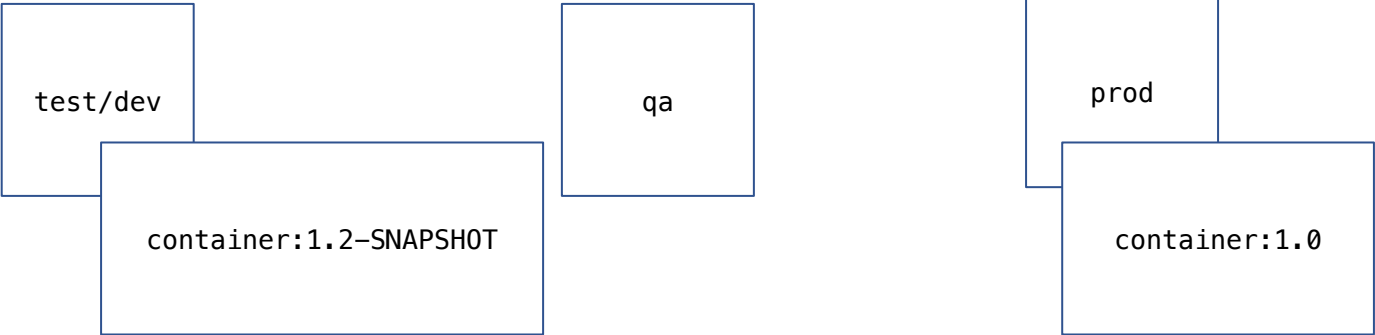
```
$ podman images
REPOSITORY          TAG          IMAGE ID      CREATED        SIZE
localhost/nginx     latest      e420c54187d7  14 seconds ago 260 MB
localhost/nginx     1          2fd45c021c45  9 minutes ago  260 MB
```

```
$ podman tag nginx:latest nginx:1
```

```
$ podman images
REPOSITORY          TAG          IMAGE ID      CREATED        SIZE
localhost/nginx     1          e420c54187d7  27 seconds ago 260 MB
localhost/nginx     latest      e420c54187d7  27 seconds ago 260 MB
<none>              <none>      2fd45c021c45  9 minutes ago  260 MB
```

```
$ podman image prune
2fd45c021c451352e18ed2383d967fd5d510d1551837446cc0f11202c7bbae05
```

```
$ podman images
REPOSITORY          TAG          IMAGE ID      CREATED        SIZE
localhost/nginx     latest      e420c54187d7  About a minute ago 260 MB
localhost/nginx     1          e420c54187d7  About a minute ago 260 MB
```



Container Registry REST-API

```
# curl -H "Authorization: Bearer $(oc whoami -t)" https://<openshift-registry>/v2/_catalog
{
  "repositories": [
    "images/nginx",
    ...
  ]
}

# curl -H "Authorization: Bearer $(oc whoami -t)" https://<openshift-registry>/v2/images/nginx/tags/list
{
  "name": "images/nginx",
  "tags": [
    "latest"
  ]
}

# curl -H "Auth ..." -H 'Accept: application/vnd.oci.image.manifest.v1+json' https://<openshift-registry>/v2/images/nginx/manifests/latest
{
  "schemaVersion": 2,
  "config": {
    "mediaType": "application/vnd.oci.image.config.v1+json",
    "digest": "sha256:54aece8b7356667b4e496c2451e818ced2d8490aa4833b681d604012968e4db2",
    "size": 468
  },
  "layers": [
    {
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
      "digest": "sha256:9cda32363c5d3aa1f39a23ef855cc98cebae6cde9b60aab904497af252d3a053",
      "size": 68902376
    }
  ]
}

# curl -H 'Accept: application/vnd.oci.image.config.v1+json'
      https://<openshift-registry>/v2/images/nginx/blobs/sha256:54aece8b7356667b4e496c2451e818ced2d8490aa4833b681d604012968e4db2
...

# curl -H 'Accept: application/vnd.oci.image.layer.v1.tar+gzip'
      https://<openshift-registry>/v2/images/nginx/blobs/sha256:9cda32363c5d3aa1f39a23ef855cc98cebae6cde9b60aab904497af252d3a053
...
```

podman save – Image Operation

erstellt ein TAR von einem Image
(Meta-Informationen, Configuration und Filesystem)

```
$ podman run -d --name ubi ubi7/ubi sleep infinity
82a21f9598b78835566487cb3e9427a9d709ef464813247693c044baa4687b2e
```

```
$ podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
82a21f9598b7	registry.access.redhat.com/ubi7/ubi:latest	sleep infinity	11 seconds ago	Up 10 seconds ago		ubi

```
$ podman images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
registry.access.redhat.com/ubi7/ubi	latest	899998a87be7	3 weeks ago	216 MB

```
$ podman save --output ubi.tar 899
```

```
$ tar -tf ubi.tar
```

```
123257361dae1cde14e6e5df3b2060adca917932129aae8a26b86c7f1e38b016.tar
c9e02f9d3afeaf029958df4ab4cdce99fc99adabc16c94975967fb5057e932c9.tar
```

```
...
```

```
repositories
```

```
manifest.json
```

```
$ podman export --output ubi-container.tar ubi
```

```
$ tar -tf ubi-container.tar
```

```
bin
```

```
boot/
```

```
dev/
```

```
etc/
```

```
etc/.pwd.lock
```

```
etc/DIR_COLORS
```

```
...
```

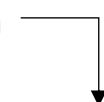
podman export – Container Operation

erstellt ein TAR von einem Container – Filesystem
ohne Meta-Information und Configuration

Container – Image

```
$ podman save <image> | tar -xf -  
$ tree -L 1 .
```

```
.  
├── 7076fcda2bf4ccbf058c10666d4c9dc2b4b643d3b6f770ed328c505387d21360  
├── 7076fcda2bf4ccbf058c10666d4c9dc2b4b643d3b6f770ed328c505387d21360.tar  
├── cbadeb4613603e1251cd6a24d6f2aa1d1bcd14a4fd2b85375e97d72a7a22764b.json  
├── manifest.json  
└── repositories
```



```
[  
  {  
    "Config": "cbadeb4613603e1251cd6a24d6f2aa1d1bcd14a4fd2b85375e97d72a7a22764b.json",  
    "RepoTags": [  
      "localhost/nginx:latest"  
    ],  
    "Layers": [  
      "7076fcda2bf4ccbf058c10666d4c9dc2b4b643d3b6f770ed328c505387d21360.tar"  
    ]  
  }  
]
```

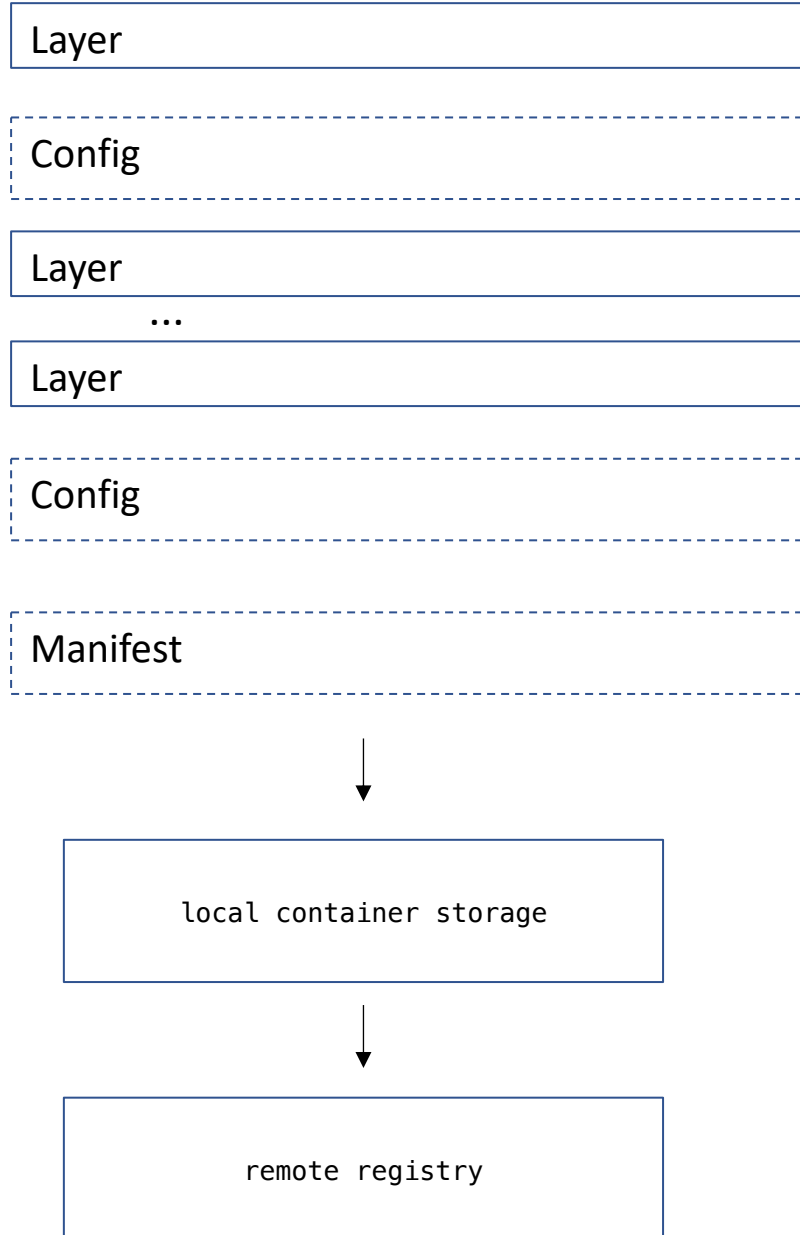
podman build - Containerfile

```
FROM ubi8/ubi
LABEL version=.... maintainer=
MAINTAINER daniel
ENV key=value
ARG version=1

ADD http://repos/app-$version.tar /opt/app/
COPY webapp.war /opt/tomcat/webapps
RUN yum install -y tomcat && \
    useradd tomcat && \
    chown tomcat:root /opt/tomcat && \
    yum cleanup && rm /tmp/* && \
    rm /var/cache/rpm

USER tomcat
WORKDIR /opt/tomcat
VOLUMES /opt/tomcat/logs
EXPOSE [ 8080, 8001 ]
ENTRYPOINT [ "bin/sh -c " ]
CMD [ "bin/catalina.sh", "start" ]
```

```
podman build -t my-tomcat .  <- Build-Dir
podman tag <registry>/<name>:<tag> <name>
podman push <registry>/<name>:<tag>
```



Verwenden von YUM/DNF beim Image-Build

```
$ sudo podman run --rm ubi8/ubi cat /etc/yum.repos.d/ubi.repo
[ubi-8-baseos]
name = Red Hat Universal Base Image 8 (RPMs) - BaseOS
baseurl = https://cdn-ubi.redhat.com/content/public/ubi/dist/ubi8/8/$basearch/baseos/os
enabled = 1
gpgkey = file:///etc/pki/rpm-gpg/RPM-GPG-KEY-redhat-release
gpgcheck = 1
...
```

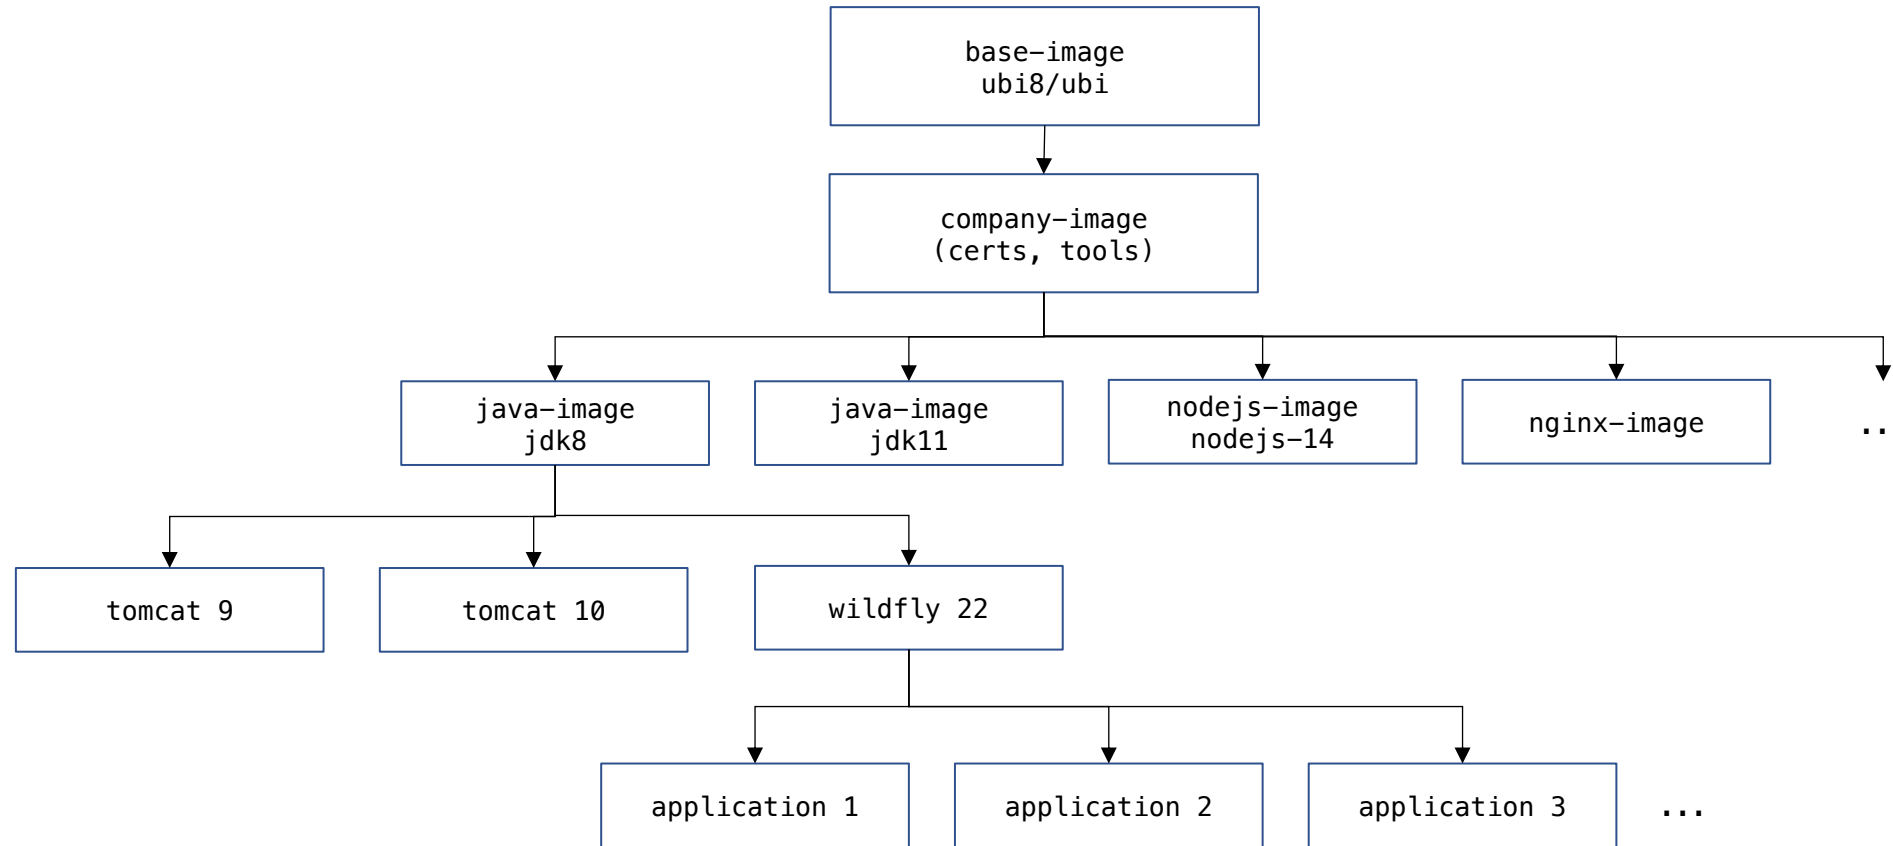
yum "telefoniert"
nach aussen !

Lösung: beim podman-build andere yum-Konfiguration (z.B. vom Host) mounten !

Bei Verwendung von Satellite/Subscriptions ggf. auch die notwendigen Zertifikate/GPG Schlüssel.

```
$ sudo podman build -v /etc/yum.repos.d:/etc/yum.repos.d -v /etc/pki:/etc/pki -v /etc/rhsm:/etc/rhsm .
```

Beispiel: Image – Vererbung



Änderungen an einem Basis-Image erfordern Rebuild der davon abhängigen Images !